



## **Project Report**

### **Artificial Intelligence – Assignment**

**Submitted to:** Cloud Counselage IT & Management Consulting Services

**Submitted by:** Mayuri Khatpe

## **Introduction**

The Artificial Intelligence Assignment focuses on building an end-to-end solution for an intelligent document classifier. The objective is to classify a set of documents into predefined categories using Natural Language Processing (NLP) and deep learning techniques. This report highlights the execution, methodology, and findings from implementing two models: Logistic Regression and LSTM.

This assignment evaluates how well insights can be derived from data, implemented through code, and communicated effectively.

## **Project Objectives**

You are tasked with building an end-to-end solution for an intelligent document classifier. The goal is to classify a set of documents into predefined categories using Natural Language Processing (NLP) and deep learning techniques.

## 1. Preprocessing Steps

The dataset used for this task consisted of 3000 text documents categorized into five classes: Scientific, Legal, E-commerce, and others. The following preprocessing steps were applied:

### Step 1.1: Data Loading

- The dataset contained three columns: `ID`, `Text` (document content), and `Category` (labels).
- No missing values were present in the dataset, and data types were consistent.
- The dataset was split into 80% training and 20% testing data.
  - Training data shape: (2400, 1685)
  - Testing data shape: (600, 1685)

### Step 1.2: Text Cleaning

- **Tokenization:** Text was split into individual words.
- **Stop-word Removal:** Common words (e.g., "the," "is") were removed using NLTK's English stop-word list.
- **Lemmatization:** Words were reduced to their base forms using the WordNet Lemmatizer.
- **Vectorization:** Text data was transformed into numerical format using **TF-IDF Vectorizer**.
  - Maximum features were limited to 5000 to reduce dimensionality.

## 2. Architecture and Methodology

Two models were implemented and evaluated for this classification task:

### A. Logistic Regression (Traditional Machine Learning)

- **Model:** A linear classifier optimized using L2 regularization.
- **Implementation:**
  - Features derived from the TF-IDF representation of text were fed into the Logistic Regression model.
  - The model was trained for 100 iterations.
- **Hyperparameter Optimization:**
  - Grid Search was used to tune `C` (regularization strength) and `solver` (optimization algorithm).
  - Best hyperparameters: `C = 100`, `solver = 'liblinear'`.
  - **Best F1-Score:** 0.9983.

### B. LSTM (Deep Learning)

- **Architecture:**
  - **Embedding Layer:** Mapped input words into a dense 128-dimensional vector space.
  - **LSTM Layer:** Captured sequential dependencies with 128 units and a 20% dropout rate.
  - **Dense Layer:** A softmax layer for multiclass classification.
- **Implementation:**
  - Features from the TF-IDF vectorizer were reshaped to fit the LSTM's input format.
  - The model was trained for 5 epochs with a batch size of 32.
  - Optimizer: Adam; Loss: Sparse Categorical Crossentropy.

### 3. Evaluation Results and Comparison

Both models were evaluated using Accuracy, Precision, Recall, and F1-Score. The results are summarized below:

Metric	Logistic Regression	LSTM
Accuracy	0.995	0.218
Precision	1.00	0.047
Recall	0.995	0.218
F1-Score	1.00	0.078

#### Observations:

- Logistic Regression significantly outperformed the LSTM model across all metrics, achieving near-perfect classification results.
- The LSTM model struggled, indicating a mismatch between the input representation (TF-IDF) and the sequential nature of LSTM.

## 4. Optimization

### Step 4.1: Optimize Logistic Regression (Grid Search)

Grid Search was used to systematically evaluate combinations of hyperparameters to identify the best-performing configuration for Logistic Regression.

- **Hyperparameters Tuned:**
  1. `C`: Regularization strength. Values tested: [0.01, 0.1, 1, 10, 100].
  2. `solver`: Optimization algorithms. Values tested: ['liblinear', 'lbfgs'].
- **Results:**
  - **Best Hyperparameters:** `C = 100, solver = 'liblinear'`.
  - **Best F1-Score: 0.9983.**

The optimized model was used for predictions, further solidifying the performance of Logistic Regression.

## 5. Challenges Faced and Solutions

### Challenge 1: Poor Performance of LSTM

- **Issue:** The LSTM model failed to capture meaningful patterns, resulting in low metrics.
- **Solution:** Future iterations could replace TF-IDF features with pre-trained embeddings (e.g., Word2Vec, GloVe, or BERT).

### Challenge 2: High Dimensionality of TF-IDF

- **Issue:** The large feature space increased computational overhead.
- **Solution:** Limited the number of features to the top 5000 terms.

### Challenge 3: Hyperparameter Tuning for Logistic Regression

- **Issue:** Identifying optimal hyperparameters for maximum performance.
- **Solution:** Used Grid Search with 3-fold cross-validation to systematically evaluate parameter combinations.

## 6. Conclusion

- The **Logistic Regression model** achieved outstanding performance, with an F1-Score of **1.00**, demonstrating its suitability for this classification task.
- The **LSTM model** underperformed, highlighting the importance of appropriate input representations for deep learning models.
- Future work should focus on incorporating embeddings or transformer-based architectures to enhance the deep learning model's performance.